

# IT125 SQL: MULTI- TABLE QUERIES

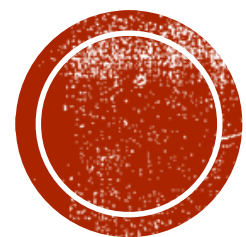
Bill Barry



# TONIGHT

- Understand the concept of a “join” and the most common types of joins
- Learn how to write SQL queries that bring data together from two tables
- Extend that knowledge to data from *more* than two tables
- Use knowledge from single-table queries; all of that applies here, too!





**JOIN CONCEPTS**



# WHAT IS A JOIN?

- Remember that exercise you did where you took data from one table and traced it through to other tables, collecting related information?
- That was a manual join you did
- You don't want to ever have to do that again, right?
- SQL will do that work for you, using SELECT's JOIN syntax

Hint: having an EER in front of you is helpful (almost *required*)





# WHAT IS A JOIN?

## MEMBER/DONATION EXAMPLE

MemberId*	MemberName	MemberPhone
1	Pat Johnson	206-555-1234
2	Lupe Valdez	206-555-4321
3	Chris Sanada	425-555-2345

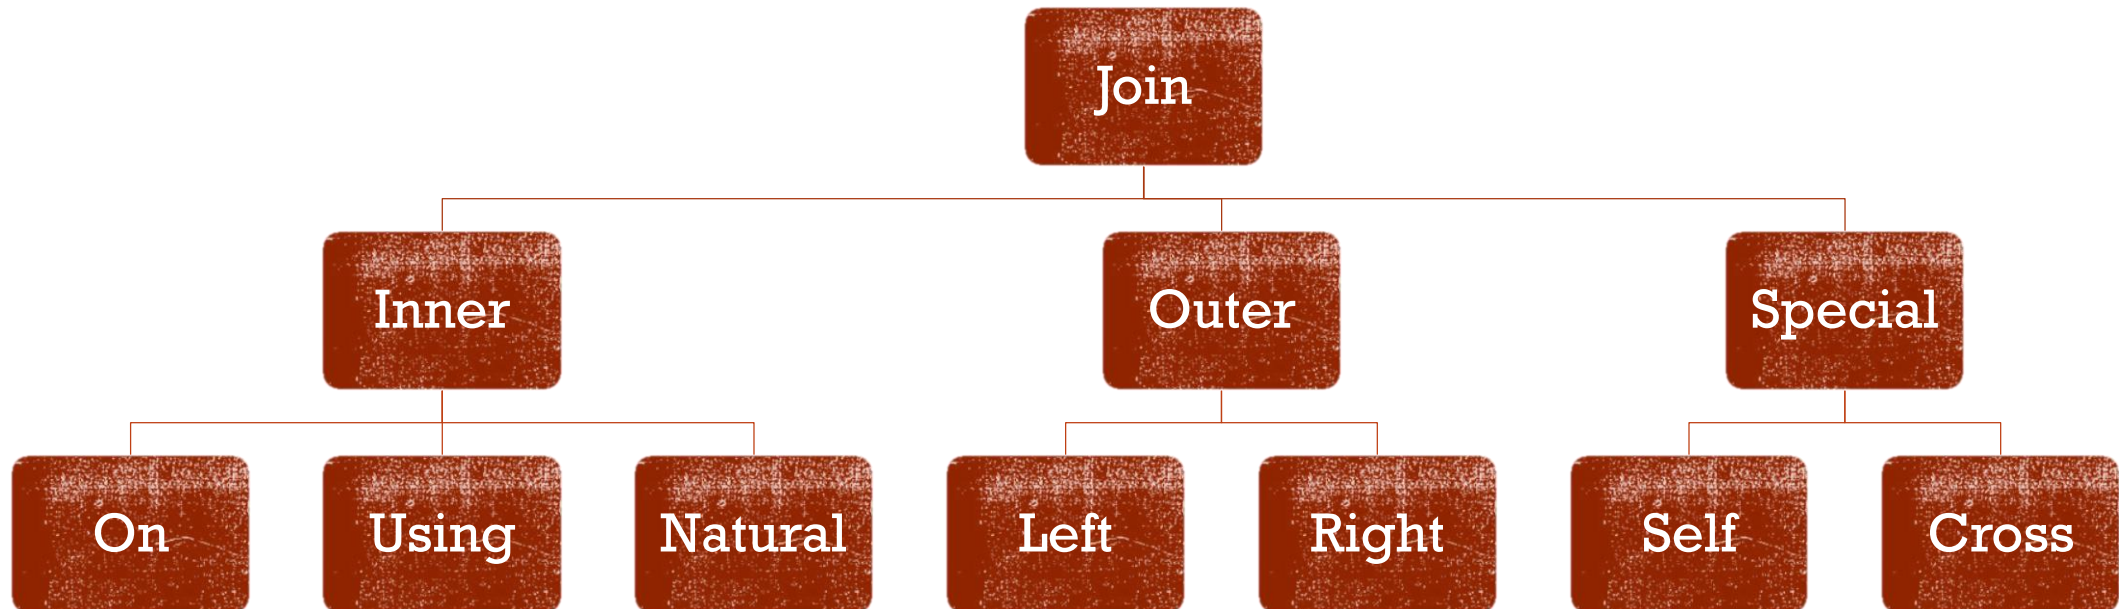
MemberId*	DonationDate*	DonationAmt
1	5/15/16	\$25
2	6/1/16	\$15
1	7/3/16	\$10
3	7/3/16	\$30

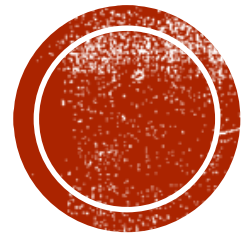


MemberId	MemberName	DonationDate	DonationAmt
1	Pat Johnson	5/15/16	\$25
1	Pat Johnson	7/3/16	\$10
2	Lupe Valdez	6/1/16	\$15
3	Chris Sanada	7/3/16	\$30



# JOIN HIERARCHY

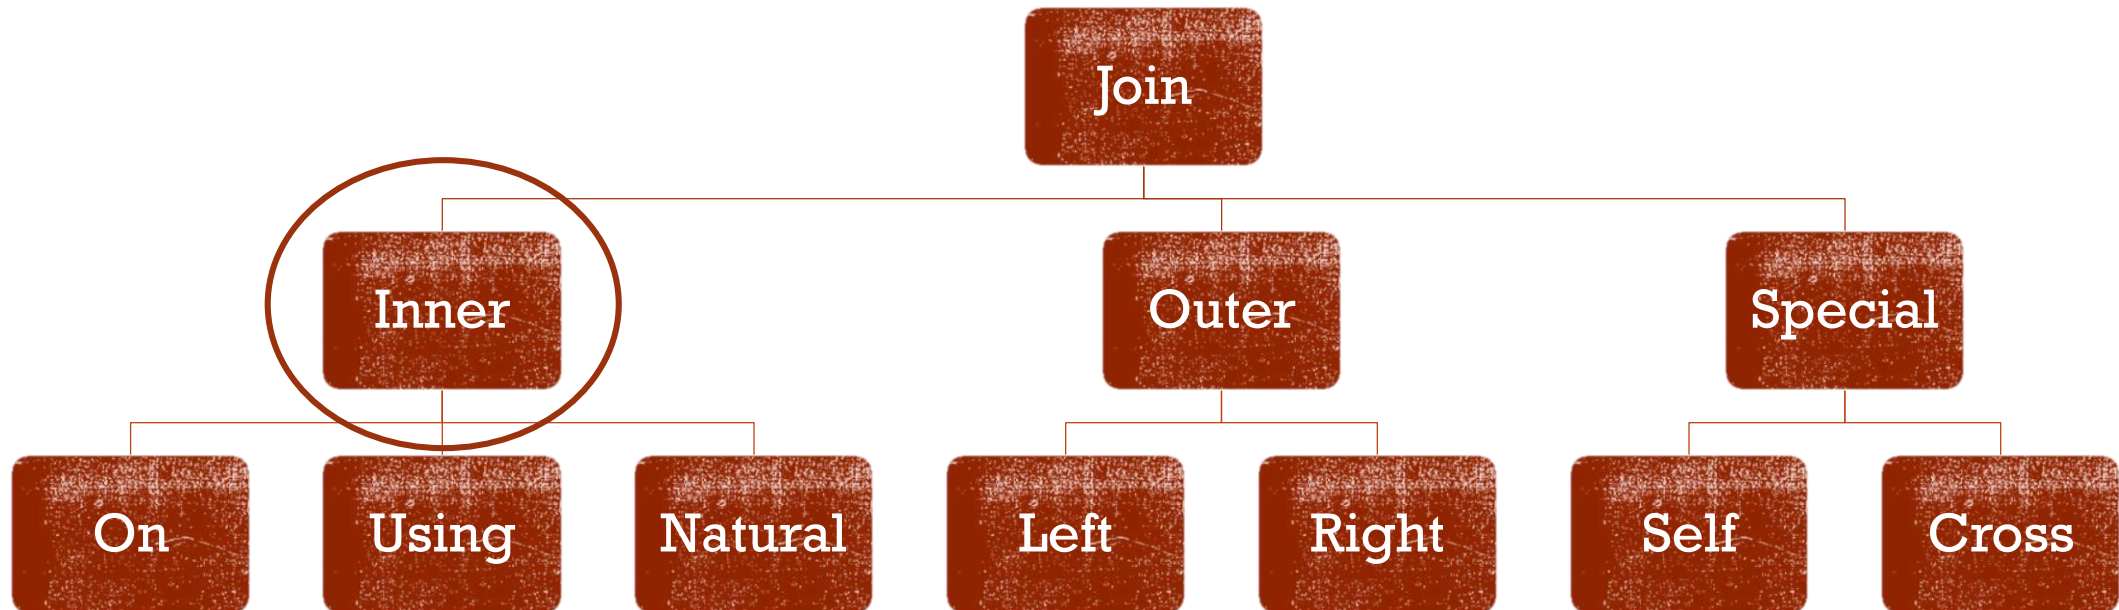




# WRITING JOINS IN SQL



# JOIN HIERARCHY





# WHAT IS AN **INNER** JOIN?

## MEMBER/DONATION EXAMPLE

MemberId*	MemberName	MemberPhone
1	Pat Johnson	206-555-1234
2	Lupe Valdez	206-555-4321
3	Chris Sanada	425-555-2345
4	Dana Smith	206-555-7654

MemberId*	DonationDate*	DonationAmt
1	5/15/16	\$25
2	6/1/16	\$15
1	7/3/16	\$10
3	7/3/16	\$30
52	6/15/16	\$50



MemberId	MemberName	DonationDate	DonationAmt
1	Pat Johnson	5/15/16	\$25
1	Pat Johnson	7/3/16	\$10
2	Lupe Valdez	6/1/16	\$15
3	Chris Sanada	7/3/16	\$30



# EXAMPLE INTRO: PET ADOPTION

- You don't have this database; we'll create it together in a future lesson

## Customer

CustId	CustName	CustPhone	CustBalance
1	Judd Jetson	2068881414	34.45
2	Harriett Hanson	4258882626	0.00
3	Betty Beaumont	2068884747	75.00

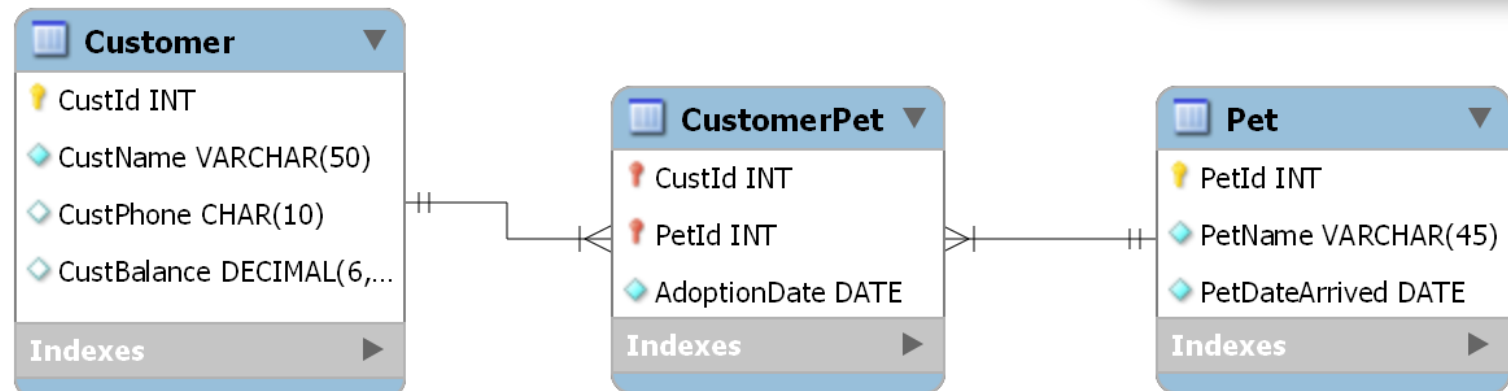
## CustomerPet

CustId	PetId	AdoptionDate
2	2	2017-11-03
3	1	2017-11-06

## Pet

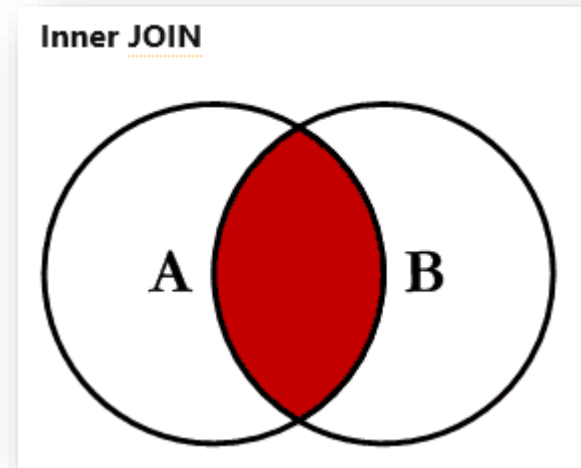
PetId	PetName	PetDateArrived
1	Fluffy	2017-10-13
2	Lucky	2017-10-07
3	Duke	2017-10-31

Remember that what makes JOINS possible is related *data*, not *schema*, per se. That said, look to PK/FK combinations when JOINing; supporting data will be there

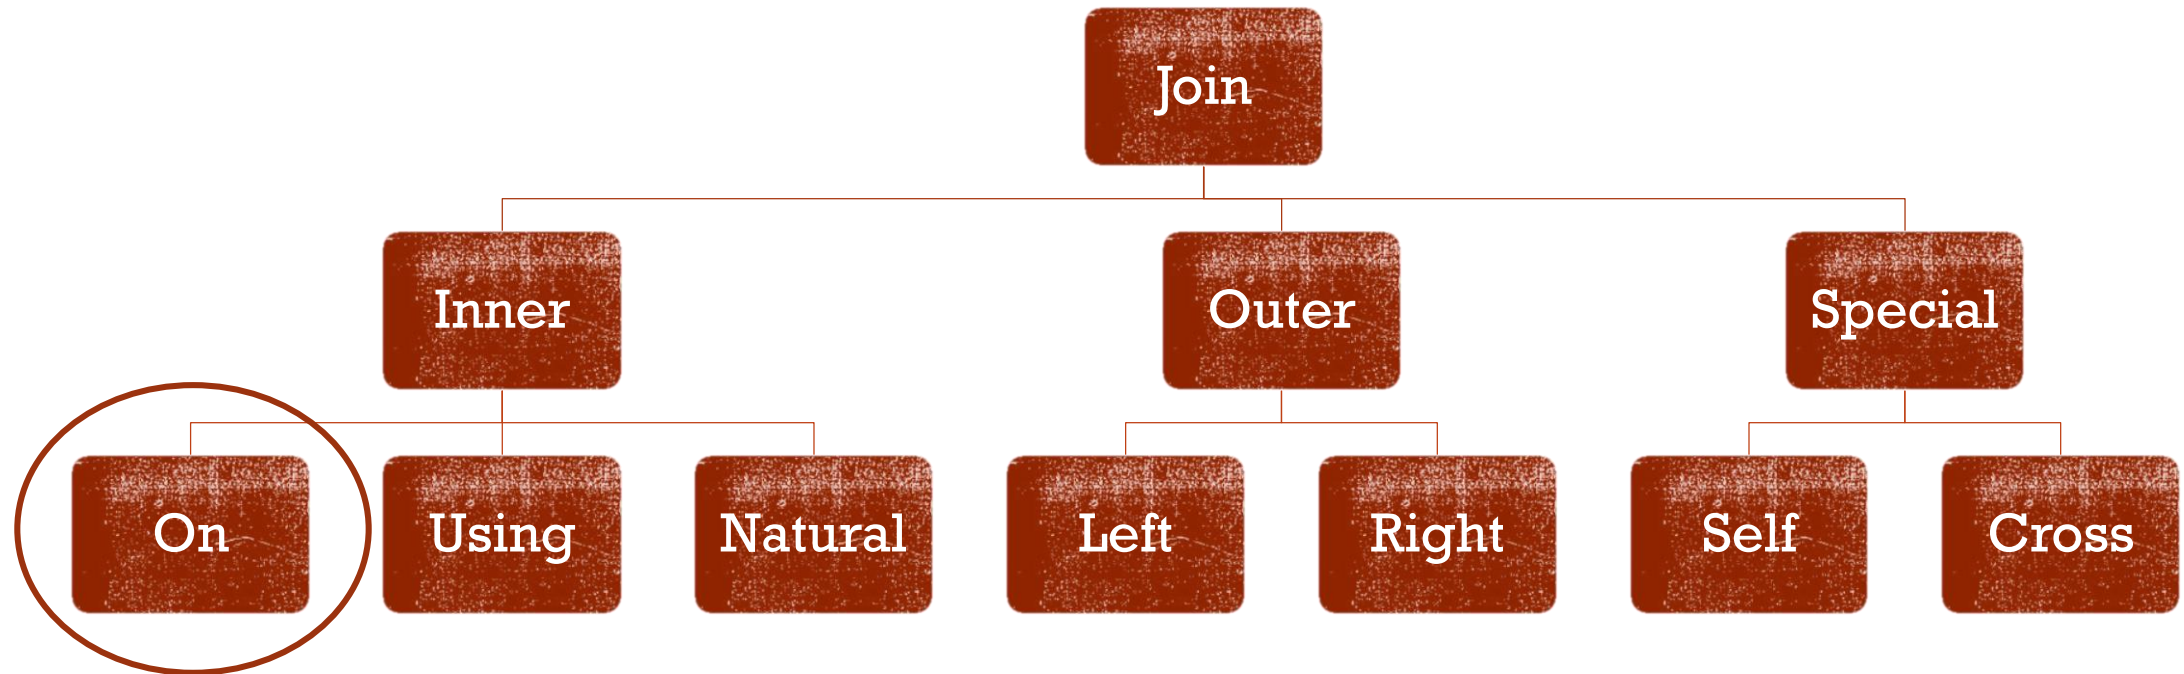


# COMMON JOIN TYPES: INNER

- The "Inner Join" is by far the most common type of join
- As shown at right, you are interested in *data that matches up* between the two tables
- An AP DB example:
  - I want a list of Invoices and their associated Vendors
  - Invoices with no associated Vendor? They won't be there!
  - Vendors who have no associated Invoices? They won't be there!



# JOIN HIERARCHY



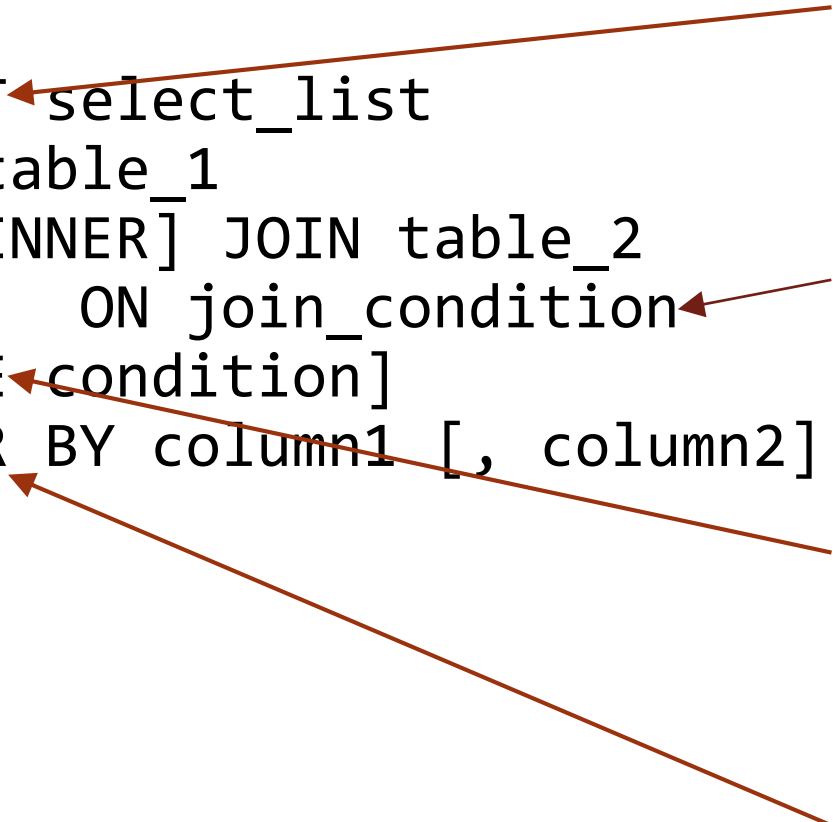
# INNER JOIN SYNTAX

- There are two major approaches:
  - Put the join condition in the FROM clause (“explicit”) ← preferred
  - Put the join condition in the WHERE clause (“implicit”)
- Let's start with the *explicit* syntax:
  - ```
SELECT select_list
  FROM table_1
      [INNER] JOIN table_2
      ON join_condition;
```
- Join conditions are typically:
  - `table1column = table2column`
- Important: if there are ambiguous column names, qualify them: `table.column`



# INNER JOIN SYNTAX ZOOM-IN

```
SELECT select_list  
FROM table_1  
    [INNER] JOIN table_2  
        ON join_condition  
[WHERE condition]  
[ORDER BY column1 [, column2]..
```



Can draw on data from either of  
the tables involved

Should focus on PK/FK columns  
that relate the two tables

Can draw on data from either of  
the tables involved

Can draw on data from either of  
the tables involved

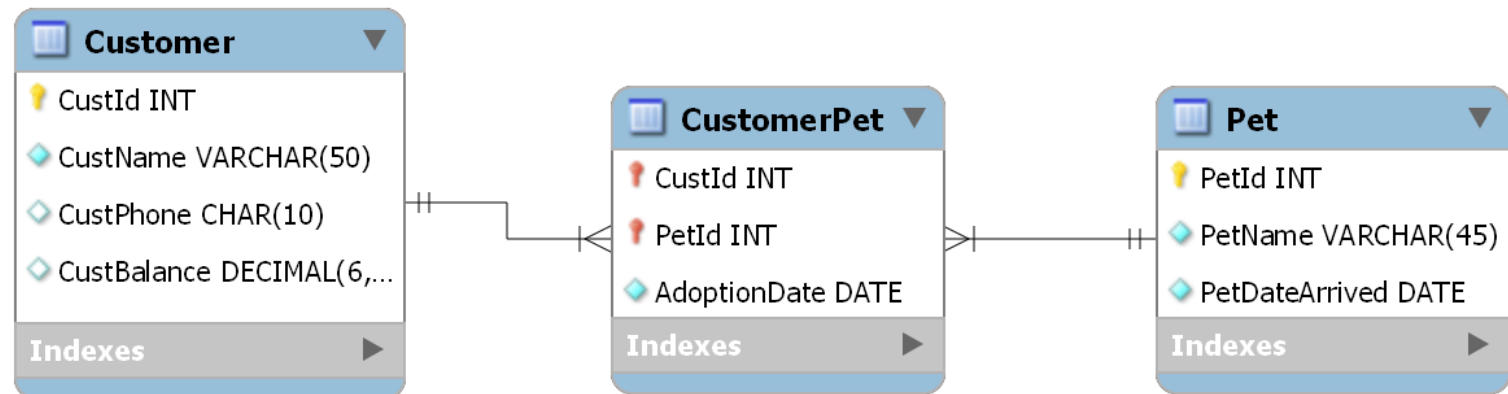


# INNER JOIN “ON” EXAMPLE

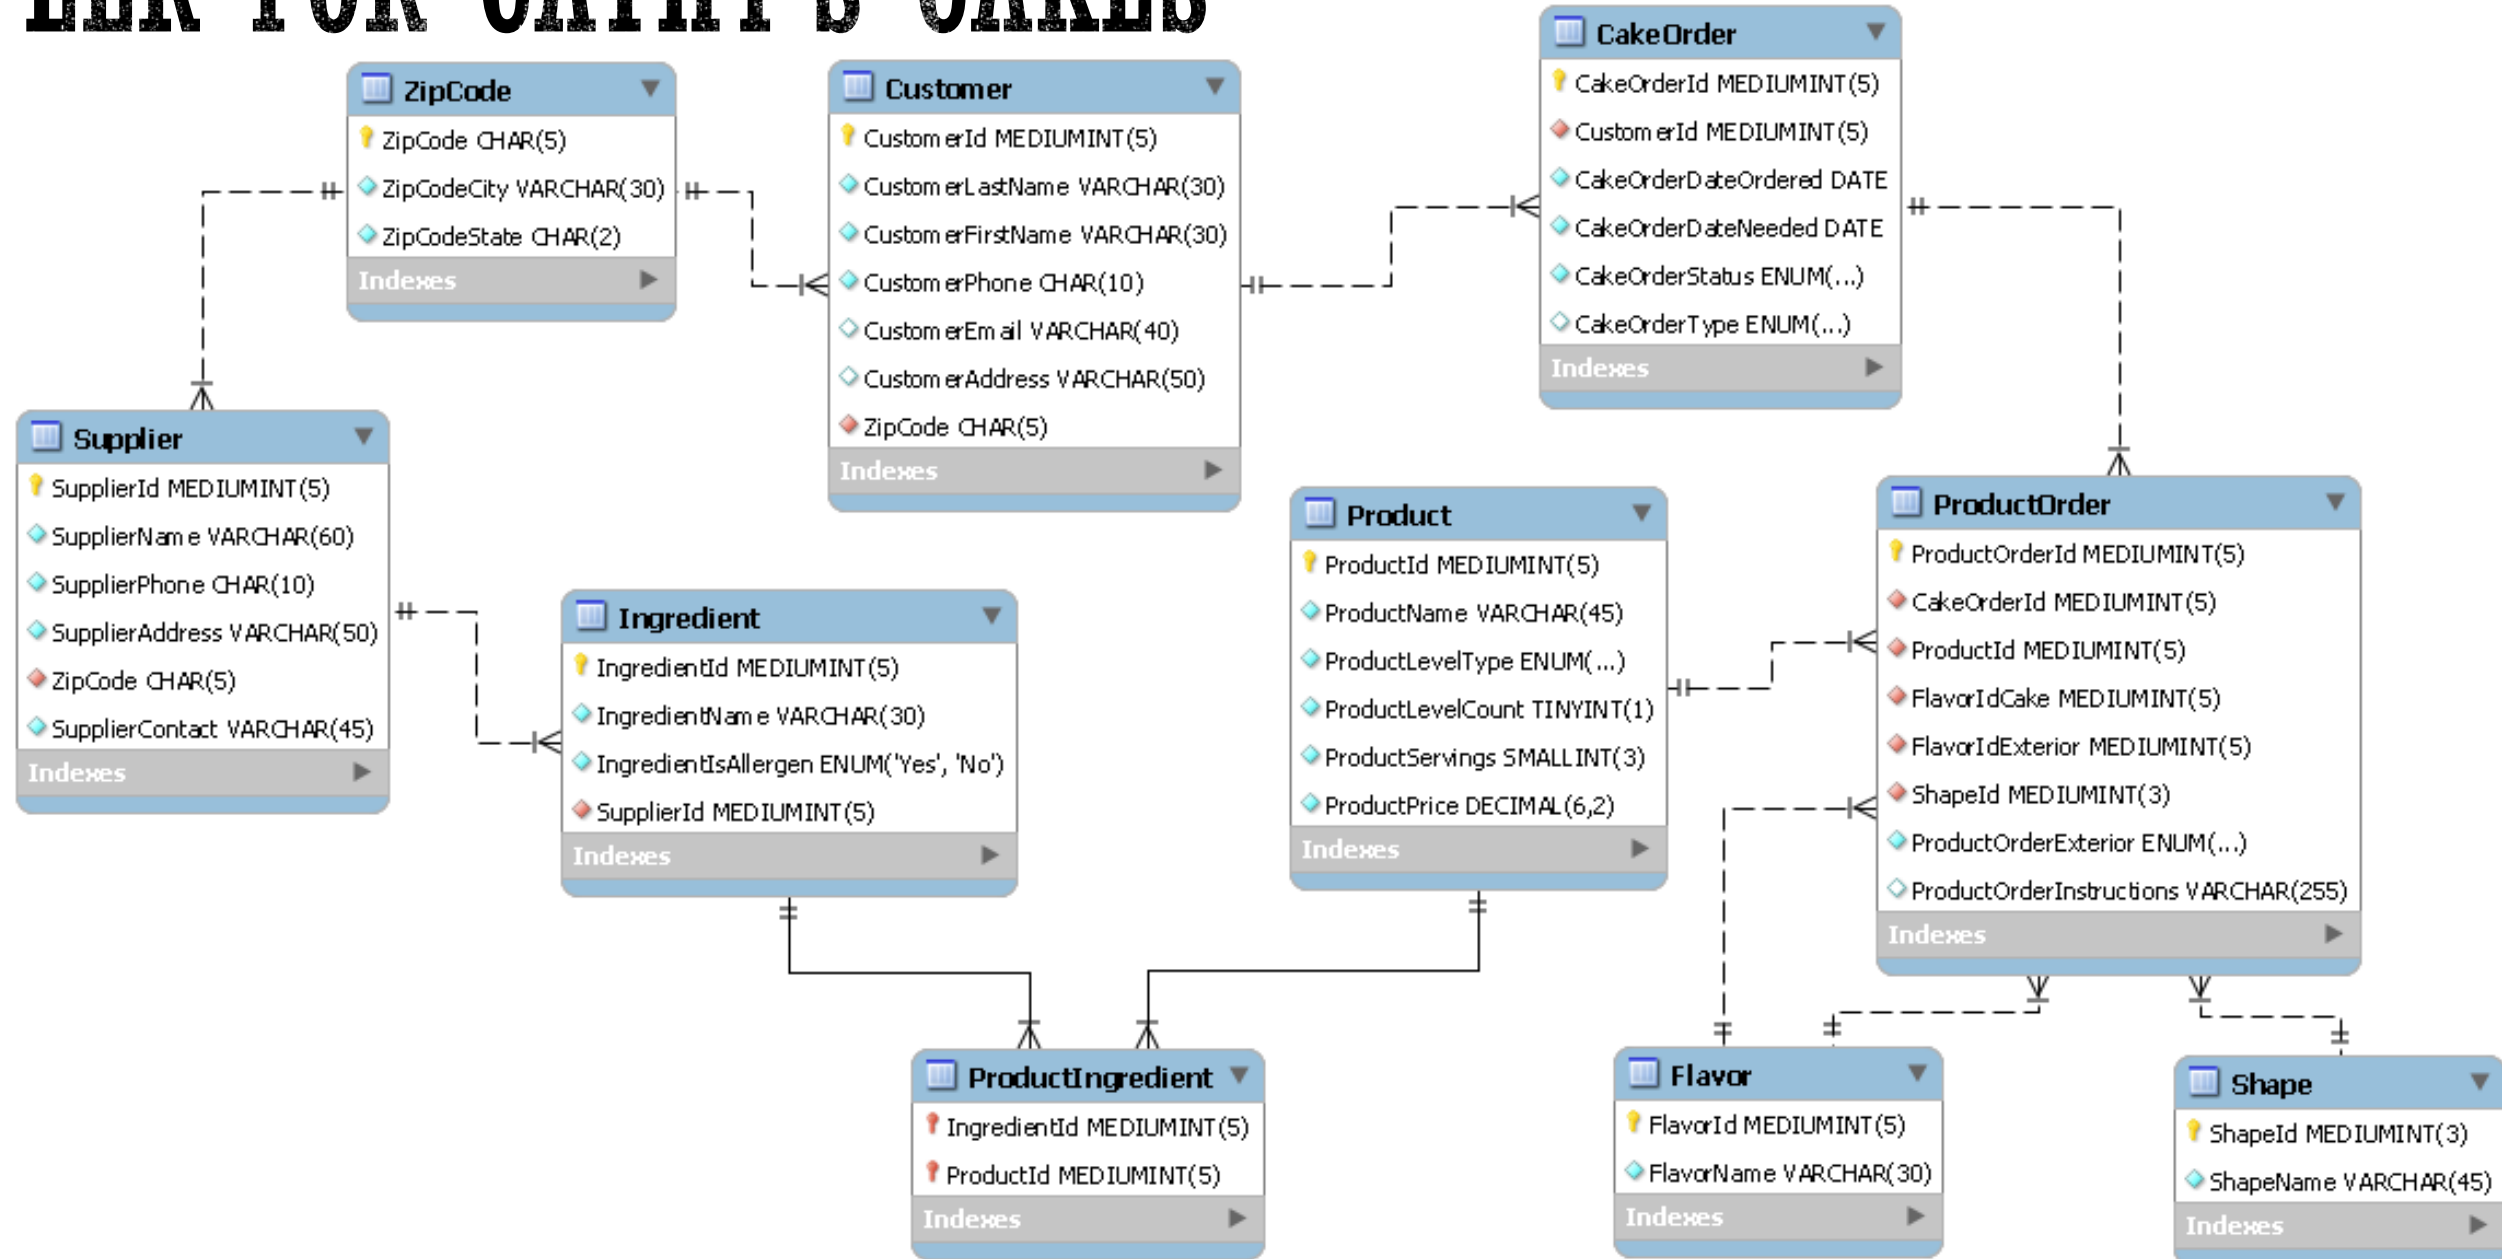
- Get a list of Pets (by name) and when they were adopted
- `SELECT PetName, AdoptionDate  
FROM Pet JOIN CustomerPet ON Pet.PetId = CustomerPet.PetId;`

| PetName | AdoptionDate |
|---------|--------------|
| Fluffy  | 2017-11-06   |
| Lucky   | 2017-11-03   |

INNER JOIN  
means that only  
*adopted* pets  
will be shown



# EER FOR CATHY'S CAKES







# INNER JOIN PRACTICE #1 (CATHY'S)

- Inner Join syntax:

- SELECT select\_list  
FROM table\_1  
[INNER] JOIN table\_2  
ON table1col = table2col

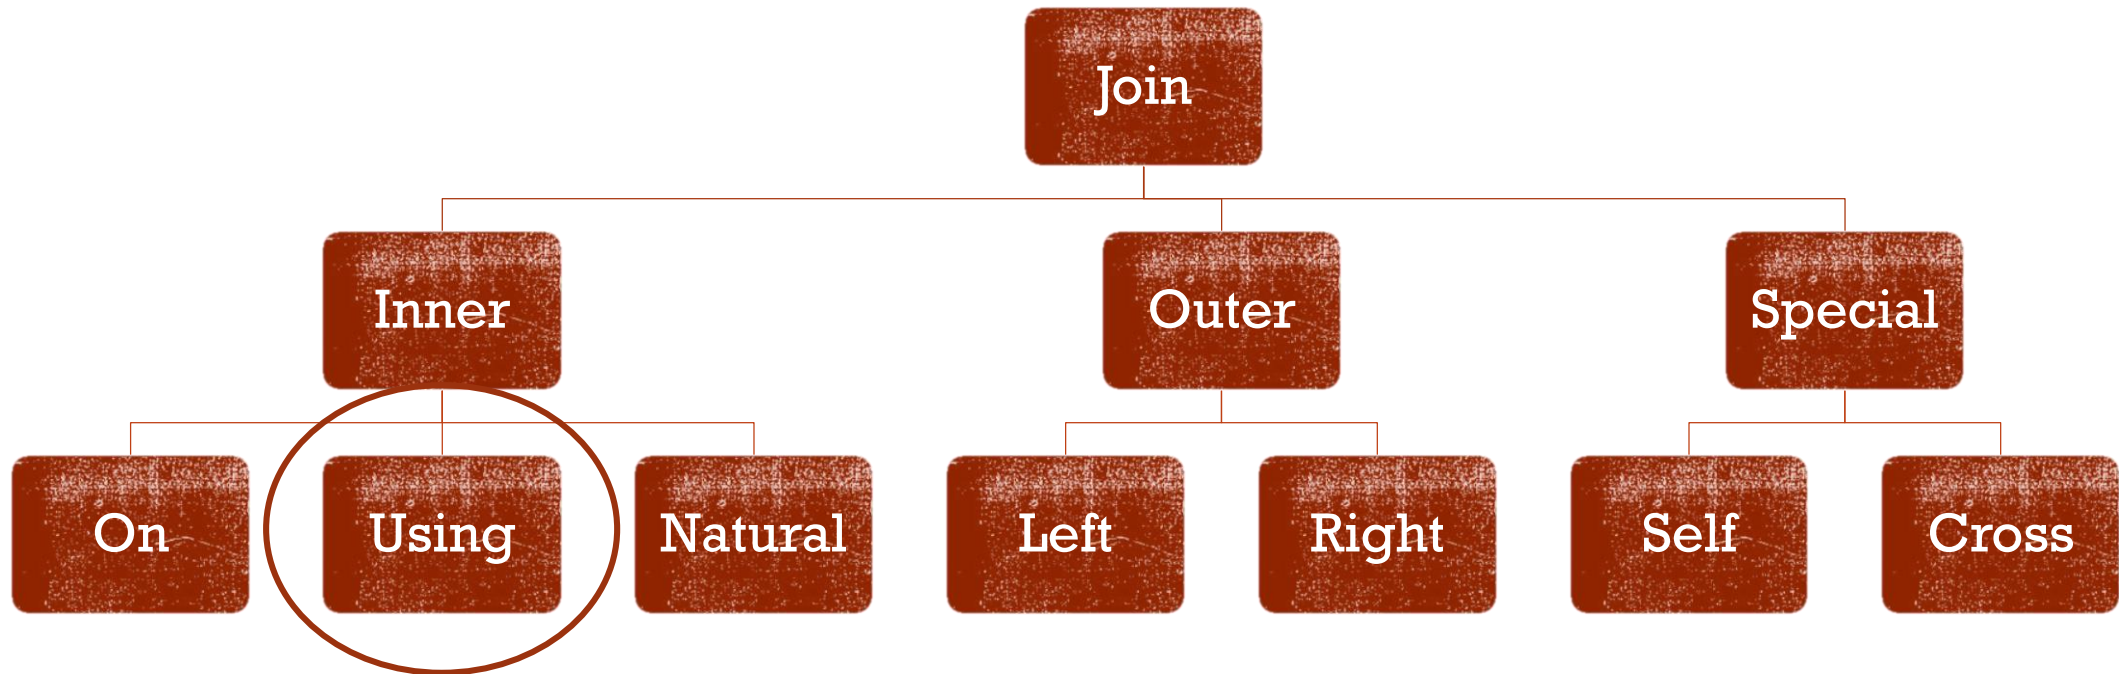
- Get a list of customers  
(full name in “last, first” order) and orders they’ve placed (id, order date, type)
- Give friendly names to each column in the results list
- Sort results by customer name, then by order date (most recent first)

| Customer          |              |
|-------------------|--------------|
| CustomerId        | MEDIUMINT(5) |
| CustomerLastName  | VARCHAR(30)  |
| CustomerFirstName | VARCHAR(30)  |
| CustomerPhone     | CHAR(10)     |
| CustomerEmail     | VARCHAR(40)  |
| CustomerAddress   | VARCHAR(50)  |
| ZipCode           | CHAR(5)      |

| CakeOrder            |              |
|----------------------|--------------|
| CakeOrderId          | MEDIUMINT(5) |
| CustomerId           | MEDIUMINT(5) |
| CakeOrderDateOrdered | DATE         |
| CakeOrderDateNeeded  | DATE         |
| CakeOrderStatus      | ENUM(...)    |
| CakeOrderType        | ENUM(...)    |
| Indexes              |              |



# JOIN HIERARCHY



# INNER JOINS: "USING" SYNTAX

- If the column name matches between the two tables, there's a shorter way:

- ```
SELECT select_list
FROM table_1
    [INNER] JOIN table_2
        USING (common_field_name);
```

- Hint: the parentheses aren't optional; you'll often forget these!

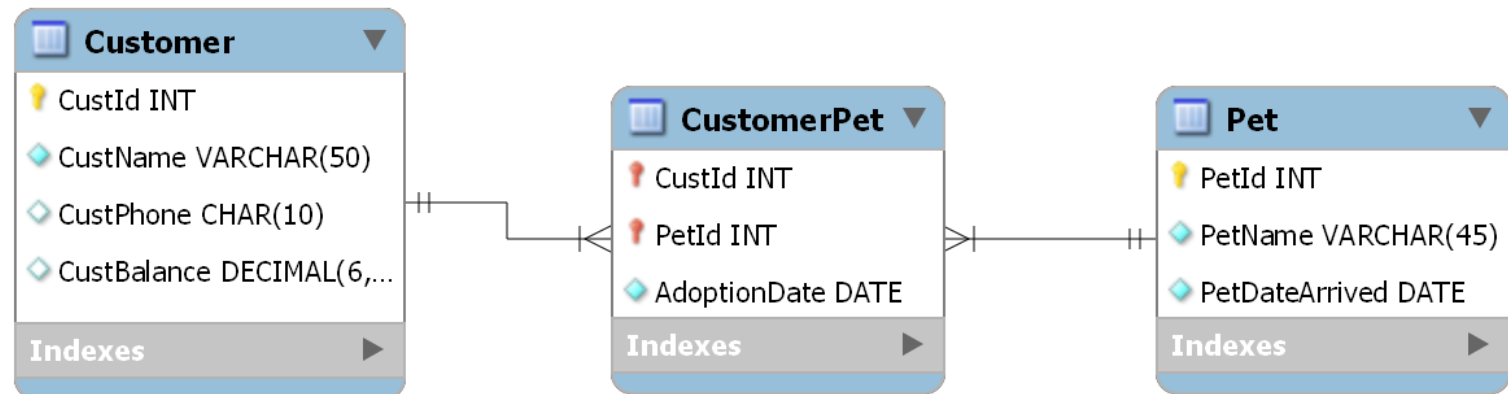


# INNER JOIN “USING” EXAMPLE

- Get a list of Pets (by name) and when they were adopted
- `SELECT PetName, AdoptionDate  
FROM Pet JOIN CustomerPet USING (PetId);`

PetName	AdoptionDate
Fluffy	2017-11-06
Lucky	2017-11-03

USING is possible here only because of the common column name

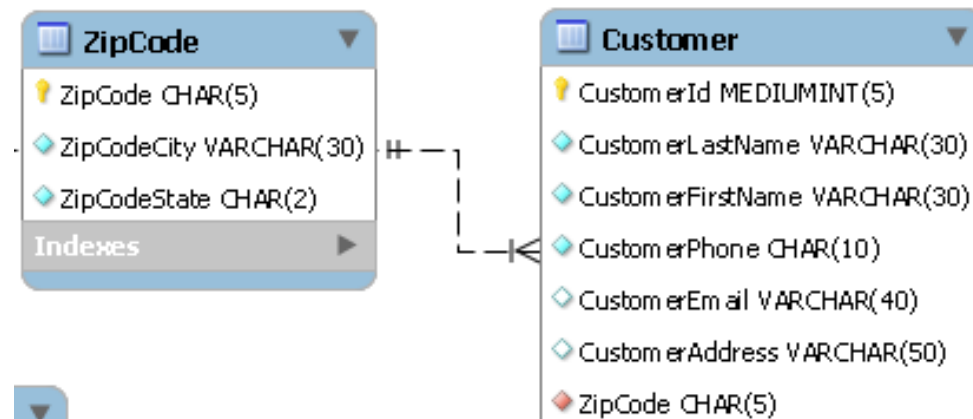




# INNER JOIN "USING" PRACTICE #2

- Inner Join "Using" syntax:

- SELECT select\_list  
FROM table\_1  
[INNER] JOIN table\_2  
USING (common\_field\_name);



- Get a list of customers who live in Bellevue, Washington
- Include the customer's name and phone number, along with the city, state, and zip
- Sort by zip code, then by customer last name, then first name
- Hints:
  - Don't look up IDs to do this; use the information provided. Same in projects!
  - If columns are ambiguous (exist in more than one table), *qualify* the column name by prepending the name of the table (e.g., *table.column*)



# TABLE ALIASES

- Sometimes it gets crazy prepending table names everywhere. SQL provides a way for you to give an alias to a table. Usually we like short ones, often one letter
- To do this, just place the alias after the table name

```
SELECT select_list
FROM table_1 [alias1]
      [INNER] JOIN table_2 [alias2]
      ON table1col = table2col;
```

Often, you'll see ALL column qualified when table aliases are used; might as well be explicit and not make the reader guess

- Then use the alias everywhere you'd prepend the table name and a dot, *even in the select list on the first line* of the query
- Example:

```
SELECT c.cust_id, c.cust_name
FROM customers c
      INNER JOIN orders o
      ON c.cust_id = o.cust_id;
```

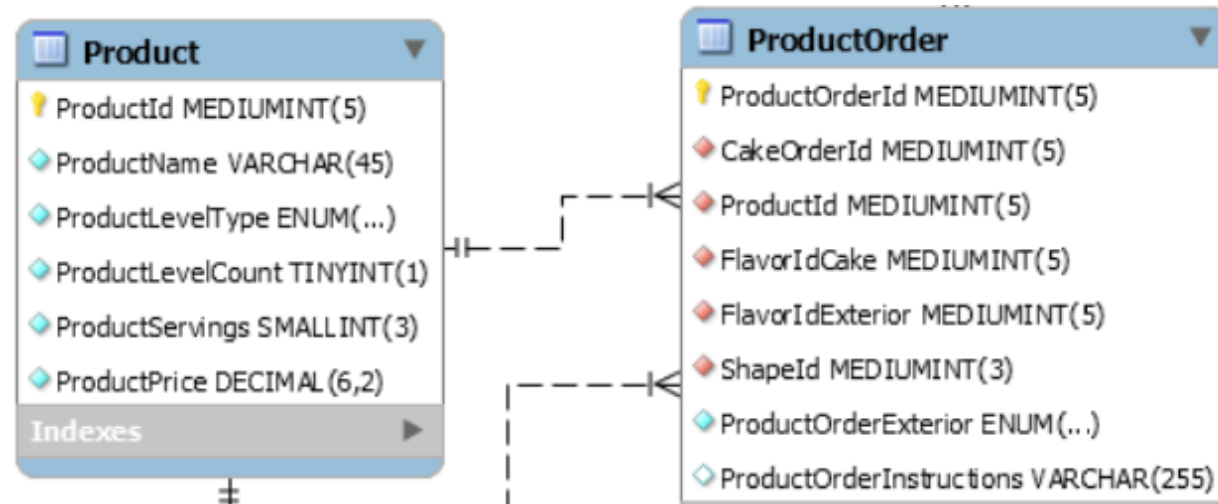
# c.cust\_name optional but clear



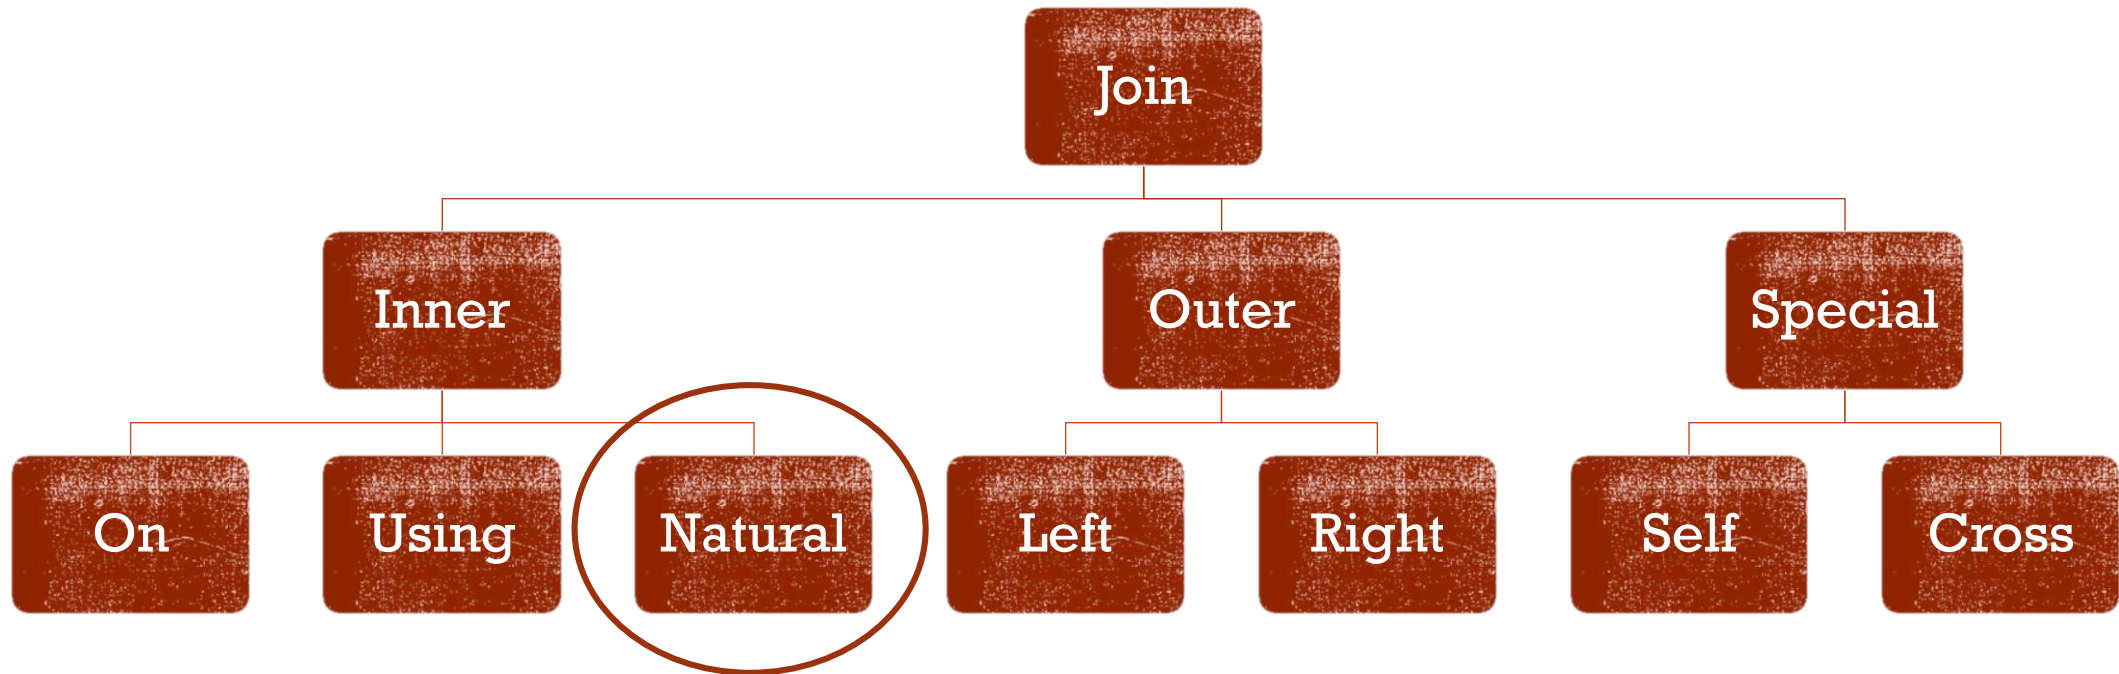


# INNER JOIN TABLE ALIASES PRACTICE #3

- Inner Join syntax:
  - `SELECT select_list`  
`FROM table_1 [alias1]`  
`[INNER] JOIN table_2 [alias2]`  
`ON table1col = table2col`
- Get a list of Order Id's along with the product ordered (by name) and prices
- Assign and use table aliases for each table
- Qualify all table-specific column references
- Sort by Order ID, then by price (highest first)



# JOIN HIERARCHY





# INNER JOINS: “NATURAL” SYNTAX

- If you trust SQL to “just do the right thing” and it can figure it out:
  - `SELECT select_list`  
`FROM table_1`  
`NATURAL JOIN table_2;`
- You get what you get and hope it's the right thing!
- In practice, this isn't generally used; we don't like the fact the results may change as the database evolves over time
- Important Caveat: “The NATURAL keyword tells the server to match up **any column names** between the two tables, and automatically use **all those columns** to resolve the join.” That means if there are two matching column names, NATURAL JOIN may fail
  - Imagine a case where each table has a TimeStamp column, for instance

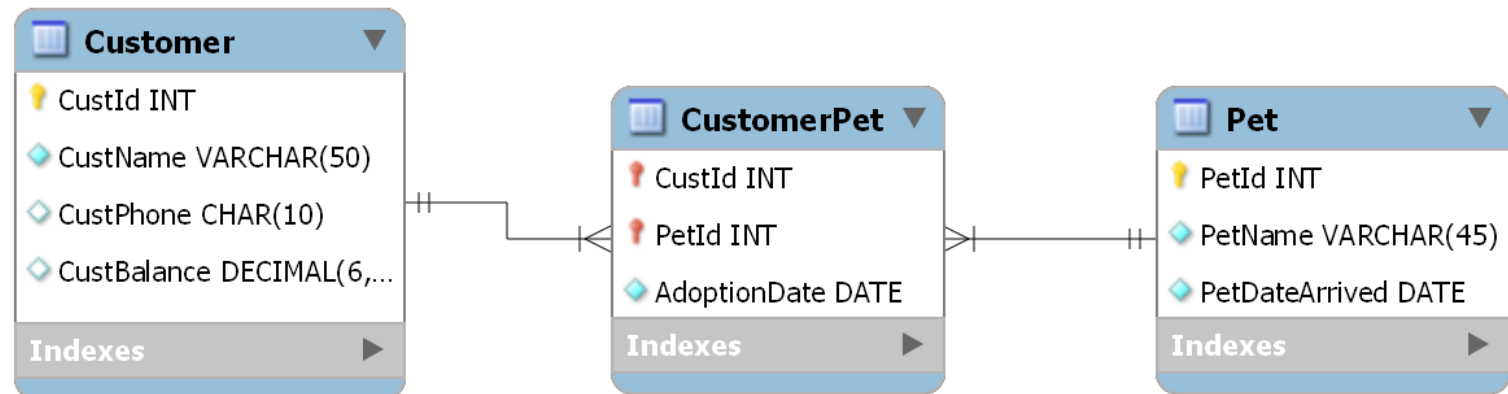


# INNER JOIN “NATURAL” EXAMPLE

- Get a list of Pets (by name) and when they were adopted
- `SELECT PetName, AdoptionDate  
FROM Pet NATURAL JOIN CustomerPet;`

PetName	AdoptionDate
Fluffy	2017-11-06
Lucky	2017-11-03

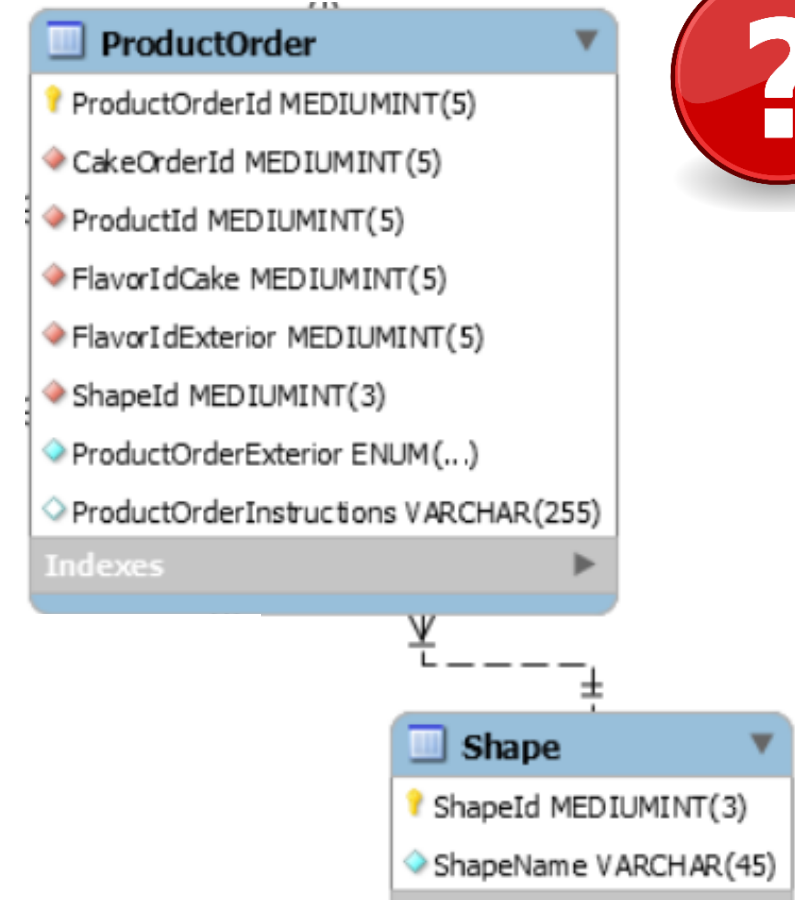
NATURAL JOIN  
works here only  
because the  
only common  
column name  
happens to be  
right JOIN  
condition



# INNER JOIN “NATURAL”

## PRACTICE #4

- Inner Join “Natural” syntax:
  - `SELECT select_list`  
`FROM table_1`  
`NATURAL JOIN table_2;`
- Find out what combinations of shapes and product exteriors have ever been ordered. Look carefully at the data to see if some additional syntax would be helpful
- Use a NATURAL JOIN to accomplish this
- Sort by exterior, then by shape
- Don't show rows where the shape name is “None” (which means there wasn't a customer choice available for this type of product, e.g., cupcakes or sheet cakes)



# INNER JOINS: IMPLICIT SYNTAX

- Another form of joins you might see is the *implicit* syntax
- This puts the join condition in the WHERE clause
  - This can make it less clear what portion of the clause is *filtering* and what portion *joining*
  - For that reason it's less preferable to the explicit syntax we've learned so far
- Syntax
  - ```
SELECT select_list  
FROM table1, table2  
WHERE table1col = table2col;
```

Note: remember to qualify (disambiguate) column names where they are the same in two tables; you may need *table1.columnName* instead of just *columnName*

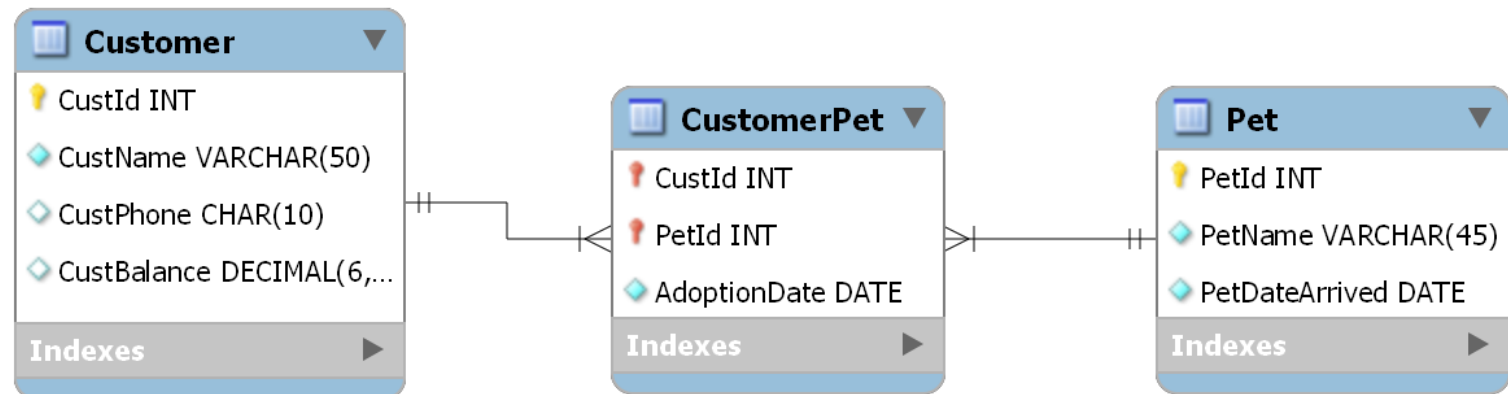


# INNER JOIN “IMPLICIT” EXAMPLE

- Get a list of Pets (by name) and when they were adopted
- ```
SELECT PetName, AdoptionDate  
FROM PetJOIN CustomerPet  
WHERE Pet.PetId = CustomerPet.PetId;
```

PetName	AdoptionDate
Fluffy	2017-11-06
Lucky	2017-11-03

Note that the JOIN condition has moved from the FROM clause to the WHERE clause

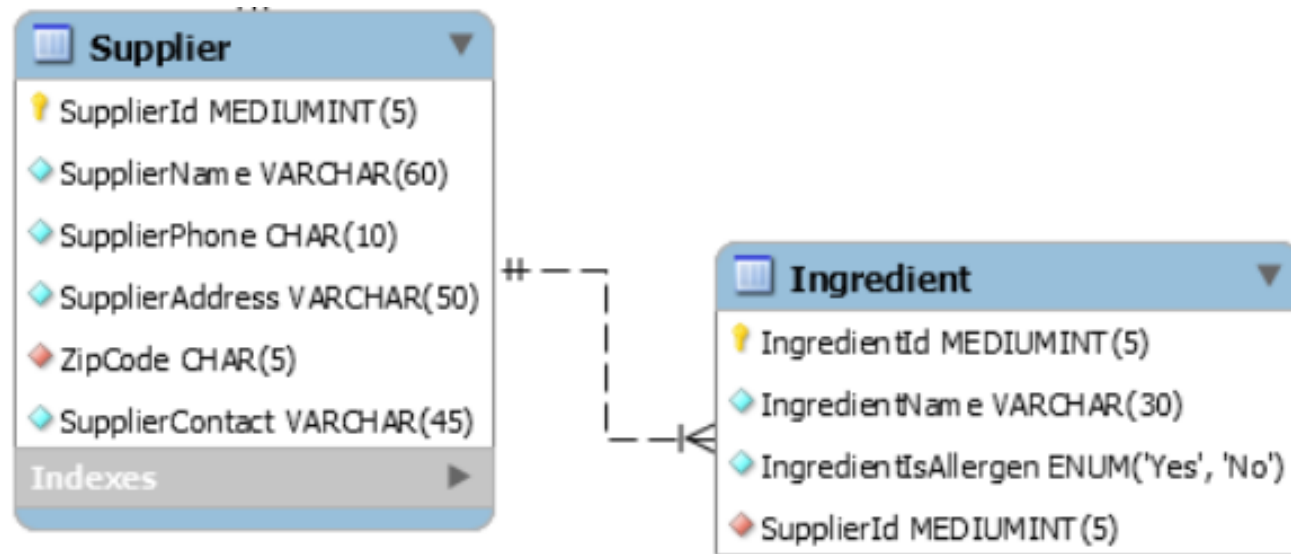




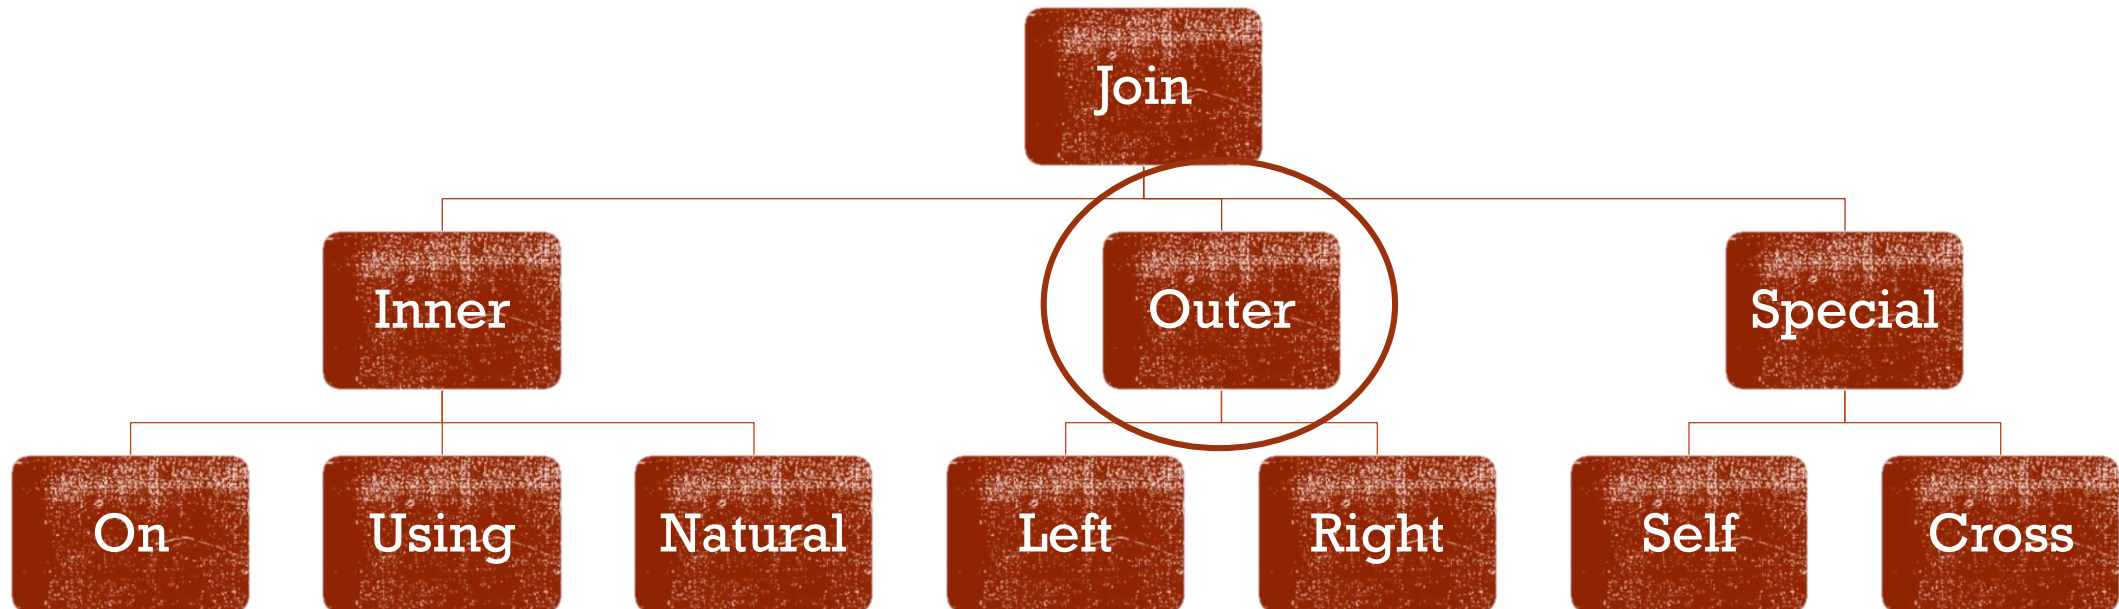
# INNER JOIN IMPLICIT SYNTAX

## PRACTICE #5

- Inner Join implicit syntax:
  - `SELECT select_list`  
`FROM table1, table2`  
`WHERE table1col = table2col;`
- Get a list of ingredients (by name) and what supplier from which Cathy buys that ingredient
- Don't include any sugar products
- Sort by ingredient name



# JOIN HIERARCHY

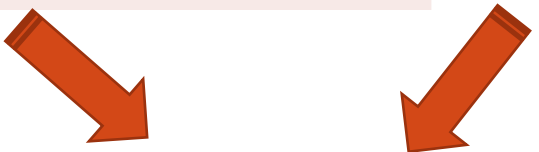


# WHAT IS AN **OUTER** JOIN?

## MEMBER/DONATION EXAMPLE

MemberId*	MemberName	MemberPhone
1	Pat Johnson	206-555-1234
2	Lupe Valdez	206-555-4321
3	Chris Sanada	425-555-2345
4	Dana Smith	206-555-7654

MemberId*	DonationDate*	DonationAmt
1	5/15/16	\$25
2	6/1/16	\$15
1	7/3/16	\$10
3	7/3/16	\$30
52	6/15/16	\$50



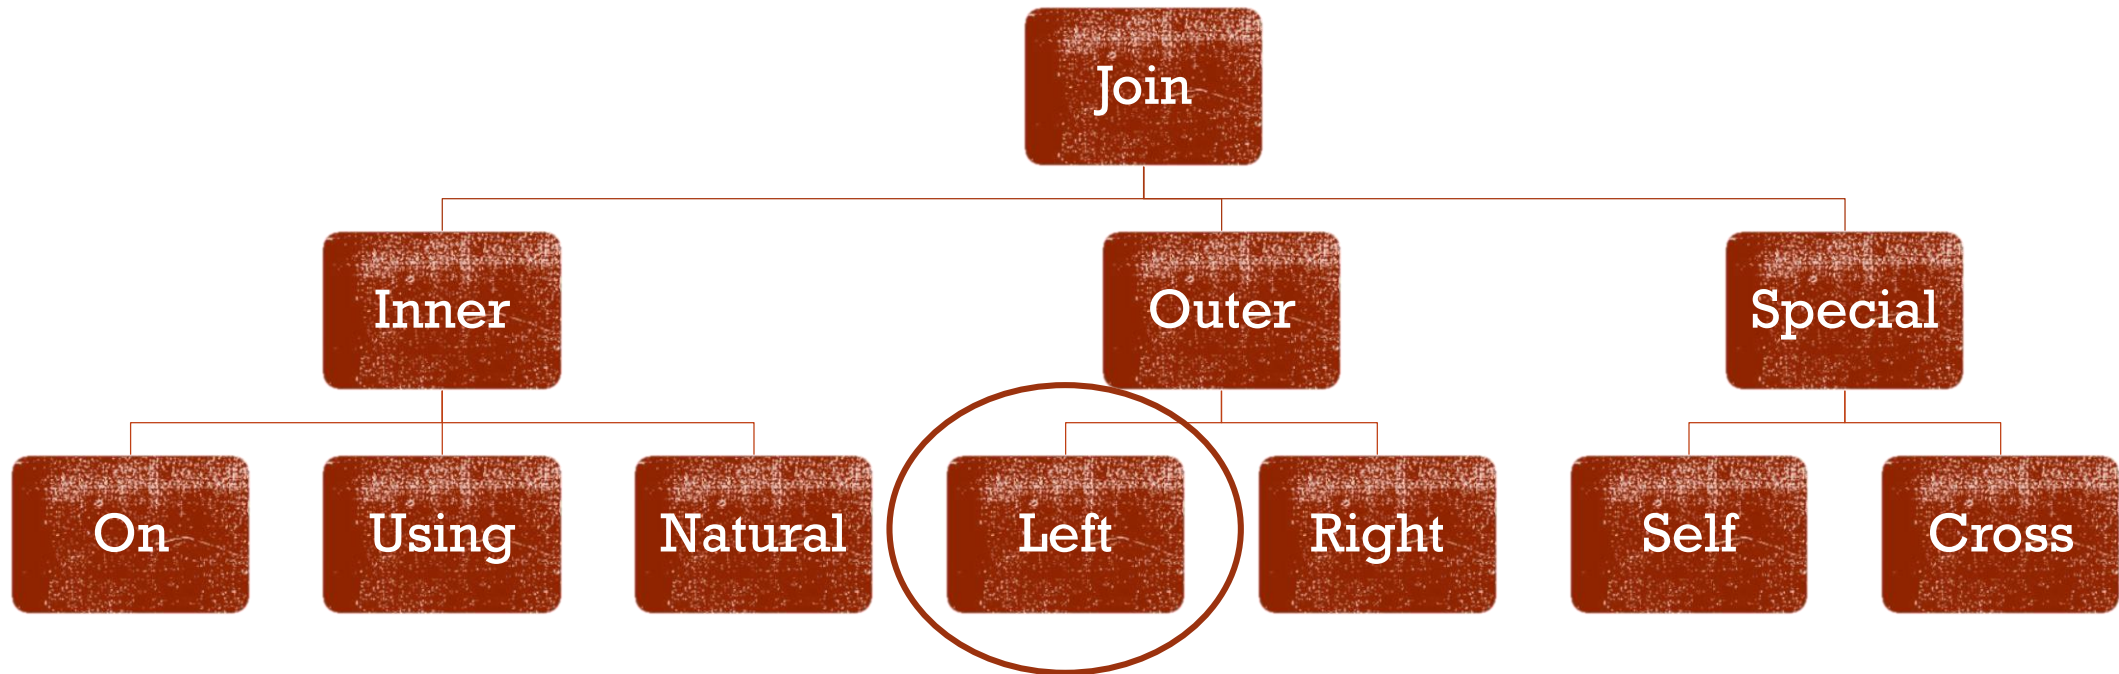
MemberId	MemberName	DonationDate	DonationAmt
1	Pat Johnson	5/15/16	\$25
1	Pat Johnson	7/3/16	\$10
2	Lupe Valdez	6/1/16	\$15
3	Chris Sanada	7/3/16	\$30
4	Dana Smith	(null)	(null)

This is a  
LEFT outer  
join—why?



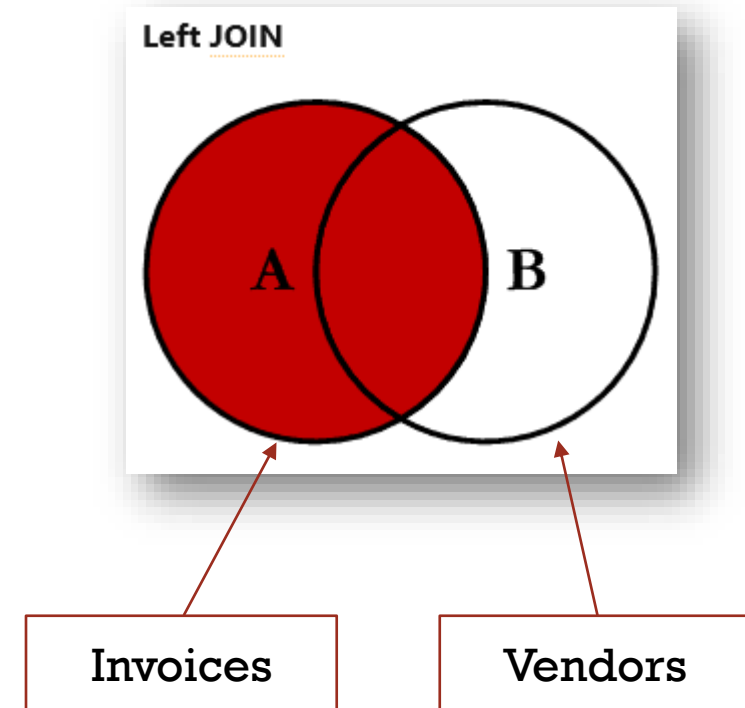


# JOIN HIERARCHY



# COMMON JOIN TYPES: LEFT OUTER JOIN

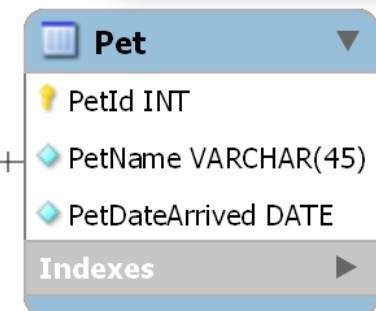
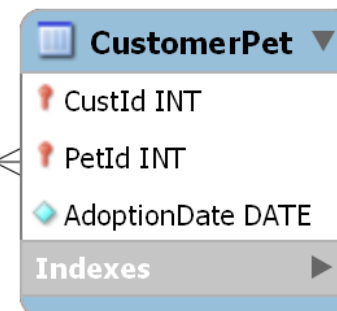
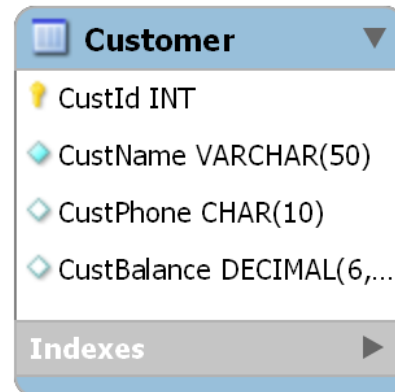
- Another common type is a Left Outer Join
- As shown at right, you are interested in all the data from the first (“left”) table, and matching data from the second table
- You shouldn't be surprised to see nulls (unmatched data)
- An AP DB example:
  - I want a list of all Invoices and associated Vendors where they exist
  - Invoices with no associated Vendor? They'll be included
  - Vendors who have no associated Invoices? Don't include them
- Syntax
  - ```
SELECT select_list
FROM table_1
    LEFT [OUTER] JOIN table_2
        ON table1col = table2col;
```



# OUTER JOIN “LEFT” EXAMPLE

- Get a list of ALL Pets (by name, regardless of their adoption status) and their associated adoption dates
- `SELECT PetName, AdoptionDate  
FROM Pet LEFT JOIN CustomerPet USING (PetId);`

| PetName | AdoptionDate |
|---------|--------------|
| Lucky   | 2017-11-03   |
| Fluffy  | 2017-11-06   |
| Duke    | NULL         |



OUTER JOINs  
*may* yield null  
data, where  
there aren't  
perfect  
matches; that's  
perfect, here

OUTER JOINs  
still require a  
JOIN condition;  
don't forget  
that!





# LEFT OUTER JOIN PRACTICE #6

- Left Outer Join syntax:

- SELECT select\_list  
FROM table\_1  
LEFT [OUTER] JOIN table\_2  
ON table1col = table2col;

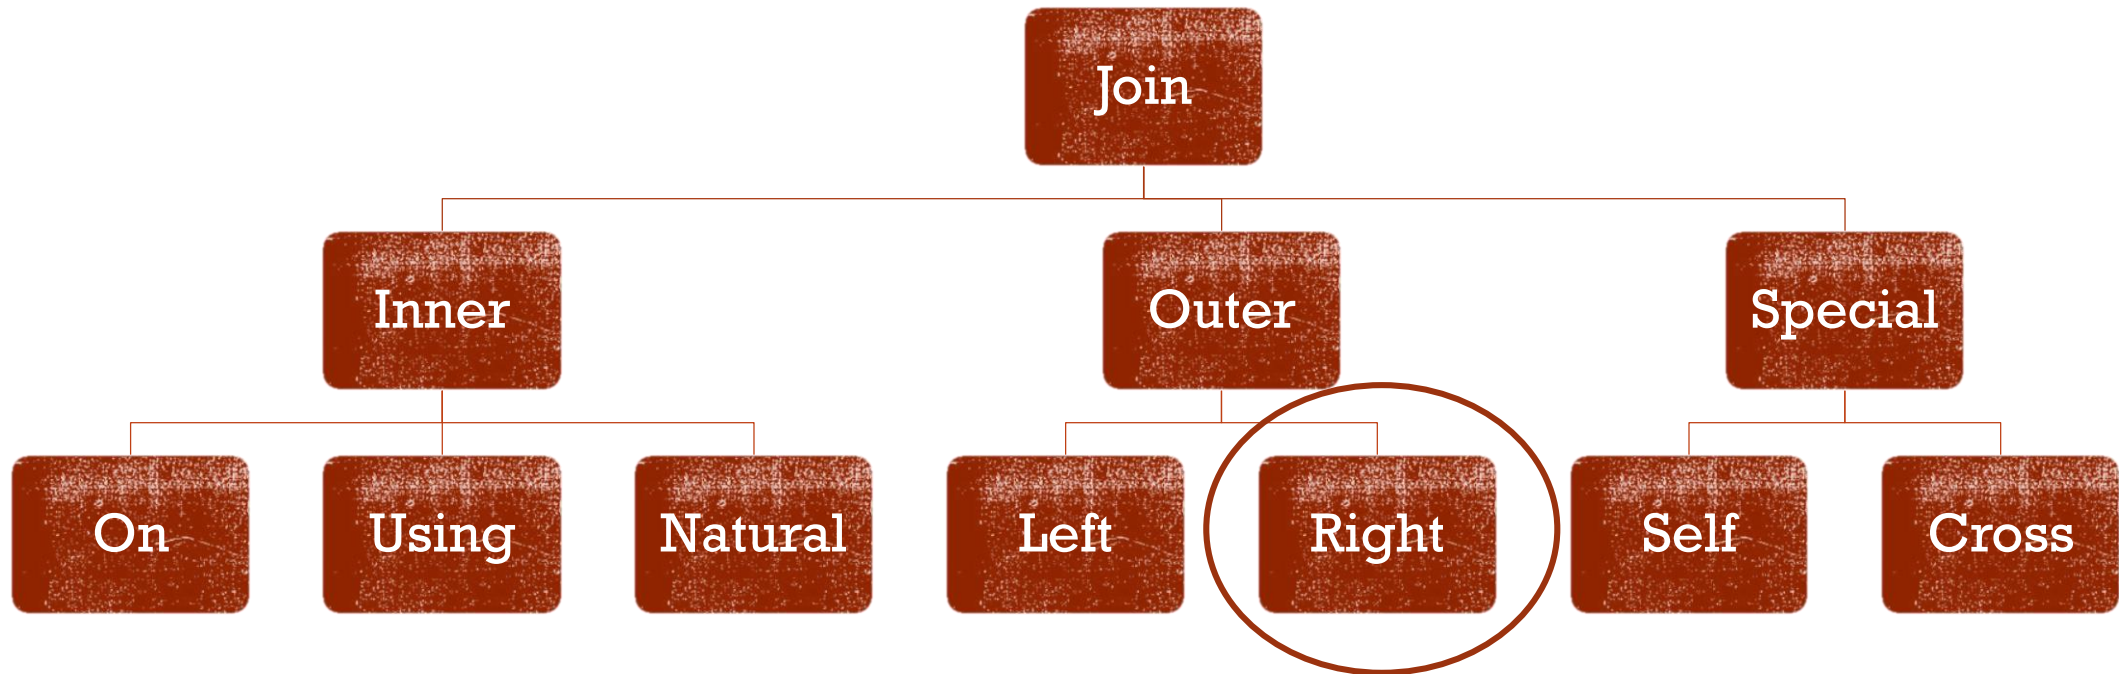
- Get a list of all customers who never placed an order
- Provide a call list that includes customer names and phone numbers
- Sort by last name, then first name
- Hints:
  - Get the join right before worrying about the “never placed an order” and the rest; develop your query incrementally
  - Don't forget that outer joins still need join conditions

| Customer          |              |
|-------------------|--------------|
| CustomerId        | MEDIUMINT(5) |
| CustomerLastName  | VARCHAR(30)  |
| CustomerFirstName | VARCHAR(30)  |
| CustomerPhone     | CHAR(10)     |
| CustomerEmail     | VARCHAR(40)  |
| CustomerAddress   | VARCHAR(50)  |
| ZipCode           | CHAR(5)      |

| CakeOrder            |              |
|----------------------|--------------|
| CakeOrderId          | MEDIUMINT(5) |
| CustomerId           | MEDIUMINT(5) |
| CakeOrderDateOrdered | DATE         |
| CakeOrderDateNeeded  | DATE         |
| CakeOrderStatus      | ENUM(...)    |
| CakeOrderType        | ENUM(...)    |
| Indexes              |              |

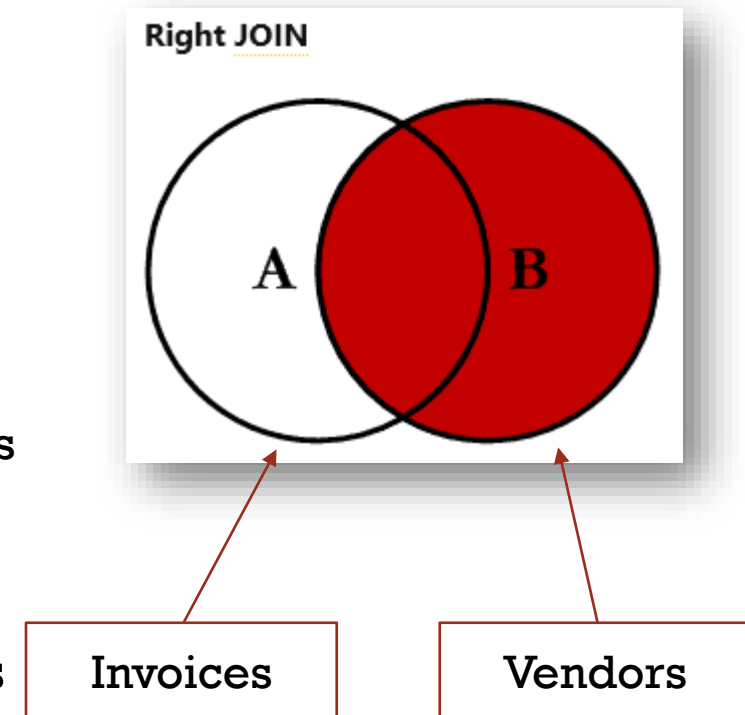


# JOIN HIERARCHY

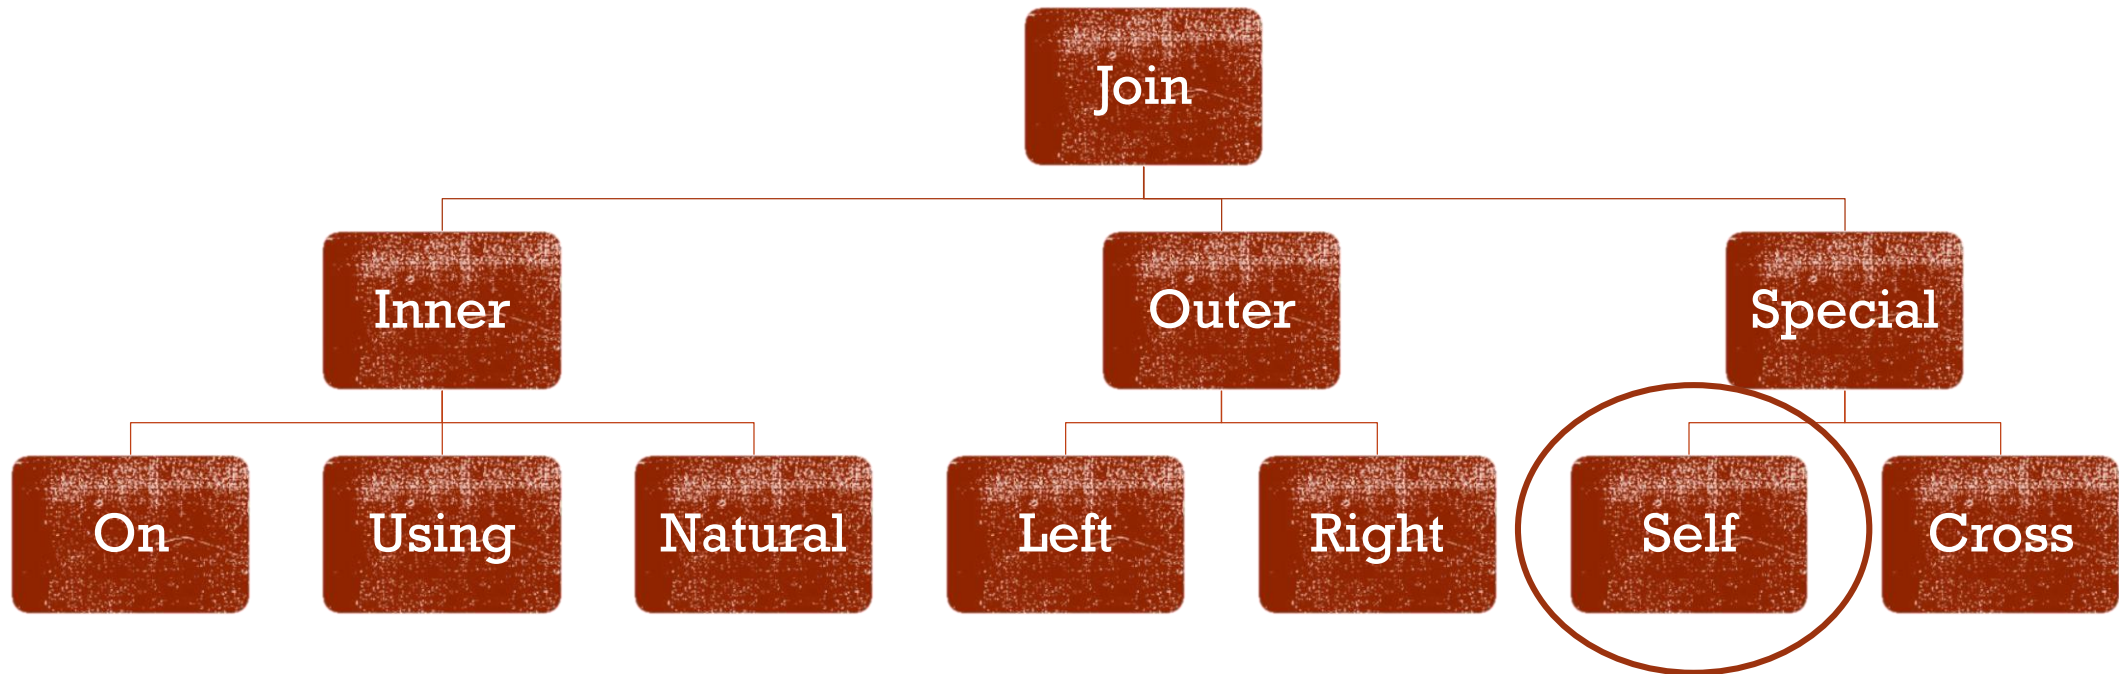


# COMMON JOIN TYPES: RIGHT OUTER JOIN

- Less common is the Right Outer Join
- As shown at right, you are interested in all the data from the second (“right”) table, and matching data from the first table
- You shouldn't be surprised to see nulls (unmatched data)
- An AP DB example:
  - I want a list of matching Invoices, where they exist, but all Vendors
  - Invoices with no associated Vendor? They won't be included
  - Vendors who have no associated Invoices? They'll be included
- In practice it's recommended you do only Left Outer Joins; it's common practice to list the “all of them” table first



# JOIN HIERARCHY





# LESS COMMON JOIN TYPES: SELF JOIN

- In this join, you use aliases to inner join a single table *to itself*
- Example: from Ex's employees table, get a list of employees and their managers
  - Other interesting variations: find employees who *have no manager*, or find employees who are *not* managers

| employees                 |
|---------------------------|
| employee_id INT(11)       |
| last_name VARCHAR(35)     |
| first_name VARCHAR(35)    |
| department_number INT(11) |
| manager_id INT(11)        |

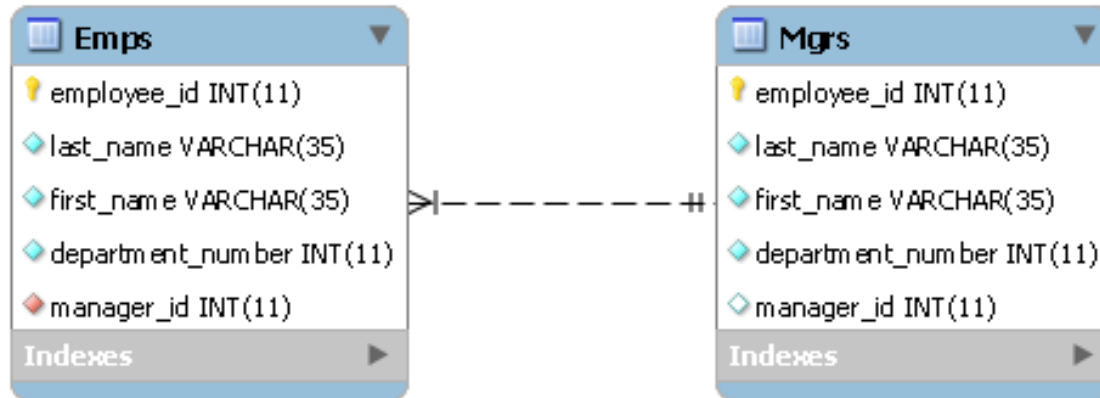
|   | employee id | last name | first name | department number | manager id |
|---|-------------|-----------|------------|-------------------|------------|
| ▶ | 1           | Smith     | Cindy      | 2                 | NULL       |
|   | 2           | Jones     | Elmer      | 4                 | 1          |
|   | 3           | Simonian  | Ralph      | 2                 | 2          |
|   | 4           | Hernandez | Olivia     | 1                 | 9          |
|   | 5           | Aaronsen  | Robert     | 2                 | 4          |
|   | 6           | Watson    | Denise     | 6                 | 8          |
|   | 7           | Hardy     | Thomas     | 5                 | 2          |
|   | 8           | O'Leary   | Rhea       | 4                 | 9          |
|   | 9           | Locario   | Paulo      | 6                 | 1          |





# HOW TO CONCEPTUALIZE A SELF-JOIN

- Our brains might explode if we try to do this directly; there's a simpler way
- Let's pretend we have **two** tables instead of **one**. Draw it this way:



- In this case the query isn't difficult; to find a list of employees and managers:  

```
SELECT Emps.last_name AS EmpLast, Emps.first_name AS EmpFirst,  
       Mgrs.last_name AS MgrLast, Mgrs.first_name AS MgrFirst  
FROM Emps JOIN Mgrs ON Emps.manager_id = Mgrs.employee_id;
```
- Now, what we we do to pretend ONE table is really TWO?
  - Anyone remember how to alias a table?



# MORPHING “PRETEND” TWO TABLE SOLUTION INTO A SELF JOIN

| EmpLast   | EmpFirst | MgrLast   | MgrFirst |
|-----------|----------|-----------|----------|
| Jones     | Elmer    | Smith     | Cindy    |
| Locario   | Paulo    | Smith     | Cindy    |
| Simonian  | Ralph    | Jones     | Elmer    |
| Hardy     | Thomas   | Jones     | Elmer    |
| Aaronsen  | Robert   | Hernandez | Olivia   |
| Watson    | Denise   | O'Leary   | Rhea     |
| Hernandez | Olivia   | Locario   | Paulo    |
| O'Leary   | Rhea     | Locario   | Paulo    |

## “Pretend” two-table SQL:

- `SELECT Emps.last_name AS EmpLast, Emps.first_name AS EmpFirst,  
Mgrs.last_name AS MgrLast, Mgrs.first_name AS MgrFirst  
FROM Emps JOIN Mgrs ON Emps.manager_id = Mgrs.employee_id;`

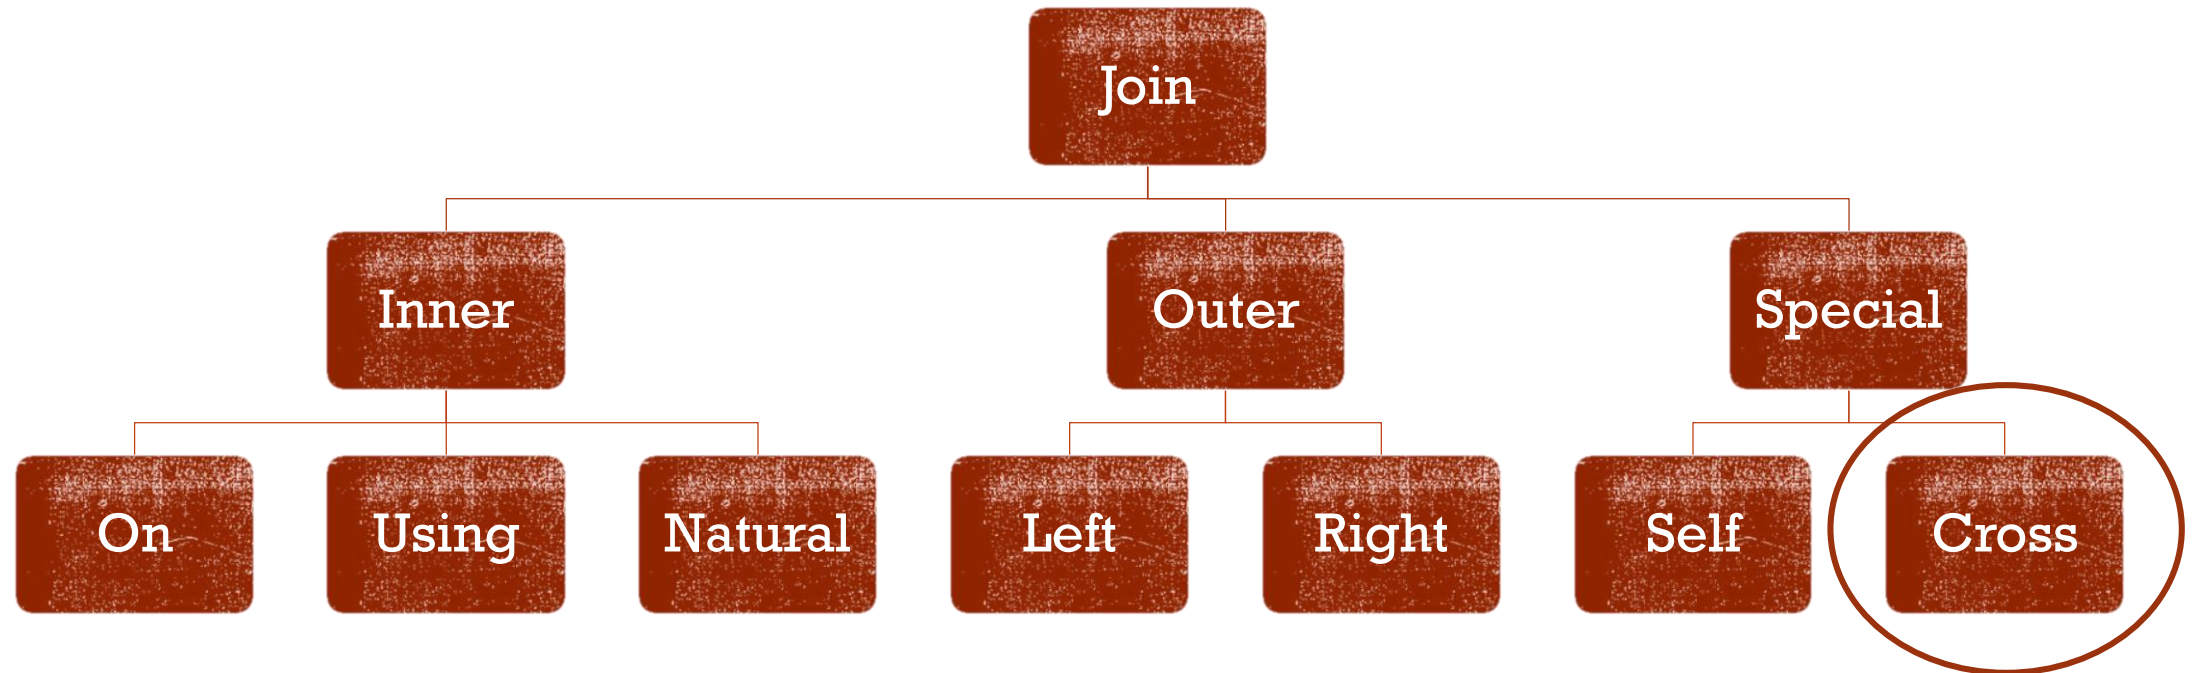
## Real Self Join SQL:

- `SELECT Emps.last_name AS EmpLast, Emps.first_name AS EmpFirst,  
Mgrs.last_name AS MgrLast, Mgrs.first_name AS MgrFirst  
FROM Employees Emps JOIN Employees Mgrs  
ON Emps.manager_id = Mgrs.employee_id;`

- Note that everything stays the same except the little aliasing trick!



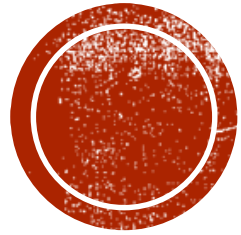
# JOIN HIERARCHY



# LESS COMMON JOIN TYPES: CROSS JOIN

- Joins *every* row in one table with *every* row in another table, a rare requirement!
- No condition is specified for this type of join
- If table1 has 20 rows and table2 has 30 rows, 600 rows will result!
- Example (ex DB): you want every department to vote on every color; get a list of all the voting that needs to be done; eliminate null color names from participation
- -- Cross join example  
SELECT department\_name AS Dept, color\_name AS Color, ' ' AS Vote  
FROM ex.departments  
CROSS JOIN ex.color\_sample  
WHERE ex.color\_sample.color\_name IS NOT NULL;
- Note: you'll get an unintentional CROSS JOIN if you forget a JOIN condition on any other type of query; that will contaminate all the data in a multi-table JOIN, too!





# JOINING MORE THAN TWO TABLES



# JOINS: NOT JUST FOR TWO!

- More tables = more interesting
- Just keep joining on more tables; start at one end, JOIN until you reach the other
- Follow the relationships (PKs and FKs), in general; again an EER is super helpful
- Typical syntax:
  - `SELECT select_list`  
`FROM table1`  
`INNER JOIN table2`  
`ON table1col = table2col`  
`INNER JOIN table3`  
`ON table2col = table3col`  
`...`
- You can use a combination of inner and outer joins if necessary; but once you use an OUTER you'll likely need OUTERS from then on to preserve the non-matching data

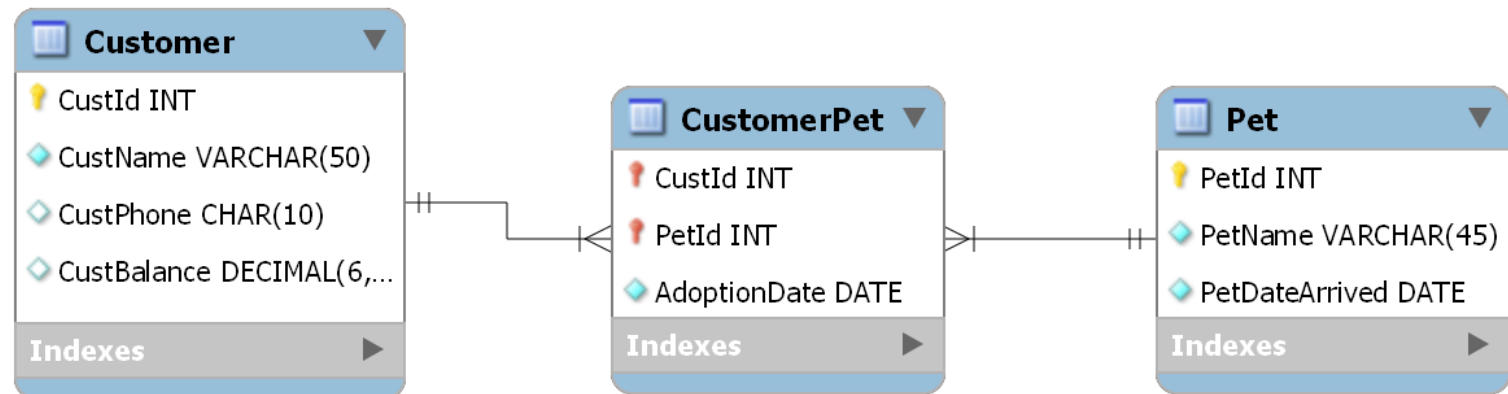


# MULTI-TABLE EXAMPLE

- Get a list of Pets (by name) and their adoptive parent's name and phone
- ```
SELECT PetName, CustName, CustPhone
FROM Pet JOIN CustomerPet USING (PetId)
      JOIN Customer      USING (CustId);
```

Don't forget that  
*each JOIN*  
needs a  
condition

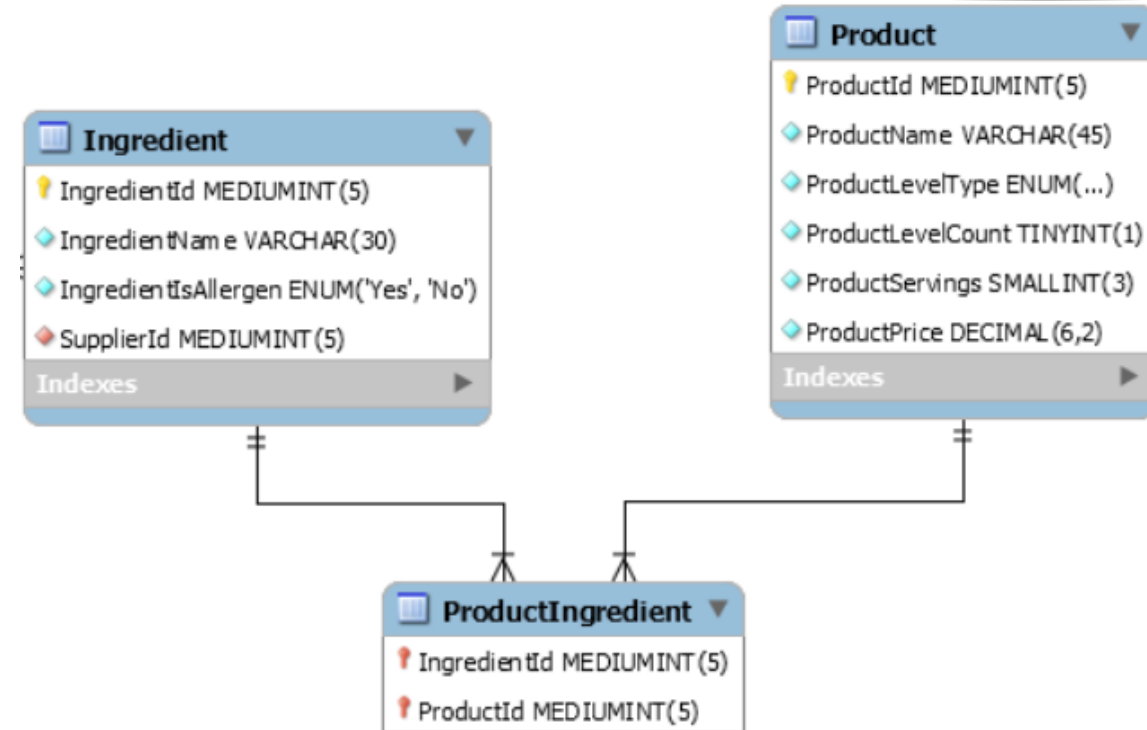
PetName	CustName	CustPhone
Lucky	Harriett Hanson	4258882626
Fluffy	Betty Beaumont	2068884747



# MULTI-TABLE PRACTICE #7



- Get a list of products along with the ingredients they contain (and whether the ingredient is a known allergen)
- Sort by ingredient, then by product name





# MULTI-TABLE PRACTICE #8



- Get a list of customers who have ever ordered any type of cupcake
- Sort by name (last, then first)
- Don't show the same customer name more than once
- Don't look up and use IDs to accomplish this (here, or on projects!)



# MULTI-TABLE PRACTICE #9



- Get a list of product orders of frosting-covered products (not fondant-covered ones) along with the cake flavor and frosting (exterior) flavor the customer chose
- Show the product name, the cake flavor name, and the exterior flavor name
- Show only cases where the two flavors aren't the same (we're looking for odd combinations). Don't repeat combinations
- Sort by product name, cake flavor name, then frosting flavor name
- Hint: write the query incrementally; this is important in complicated ones

This is a tough one; if you can do this one, you've graduated on this topic!



# TIPS FOR MULTI-TABLE QUERIES

- Figure out what tables are involved
  - Look at what data is requested; it may come from multiple tables
  - Data may not be in adjacent tables; if so, all the tables in between need to be included
- Find a path between the tables
  - Usually you can find a direct path between the tables
  - In rare cases there may be two paths; pick that one that makes the most sense given the JOIN context
- Find the join conditions
  - Look for the PK/FK combinations that relate each pair of tables; use those to construct your join condition using either ON or USING
- Figure out what kind of join is required
  - Most questions require (normal) inner joins
  - Look for wording that gives you additional clues, e.g., “find ALL customers, plus...”
  - Look out for the rarer types; you will probably see one self join and one cross join among the project questions



# WHAT SHOULD I DO NEXT?

- Take the quiz on Chapter 4
- Start Proj03
  - Querying multiple tables
  - Table design
- Before we meet again...
  - Submit Proj03
  - Read next week's material:
    - Chapter 5: Inserting, Updating, and Deleting Data
    - Chapter 8: Data Types



**QUESTIONS?**

